

Introduction to deep learning

`chyld.medford@gmail.com`

github repository

— — —

<https://github.com/chyld/intro-deep-learning>

Overview

— — —

1. Machine Learning
2. Supervised vs Unsupervised
3. Why is deep learning popular?
4. Theory of deep learning
5. Types of neural networks
6. Keras and Tensorflow
7. Code: MNIST on Dense Network
8. Theory of CNNs
9. Code: Fashion MNIST on CNN
10. Deep learning on Tesla GPUs on AWS
11. Bonus: Transfer Learning

Machine Learning


Supervised vs Unsupervised

Why is deep learning so popular?

— — —

- Lots of data
- Fast processors
- Cloud computing
- Better models

Theory of deep learning



Theory of deep learning - 1

- Perceptrons (neuron)
- Input/Output layers
- Hidden layers
- Linear Equation $w_1x_1 + w_2x_2 + w_3x_3 + \text{bias}$
- Activations (nonlinear)
 - Sigmoid, Tanh, Relu
- Feed forward
- Loss functions
- Optimizers / Backpropagation (partial derivatives)

Theory of deep learning - 2

— — —

- Update weights, biases
- Stochastic, Mini-batch, Batch update
- Epochs
- Output: Linear, Sigmoid, Softmax
- Train, validation, test
- Bias vs variance

Types of neural networks

- Dense / Fully Connected / Multilayer Perceptron
- Convolutional Neural Networks (CNN)
- Recurrent Neural Networks (RNN)
- Generative adversarial network (GAN)
- more...


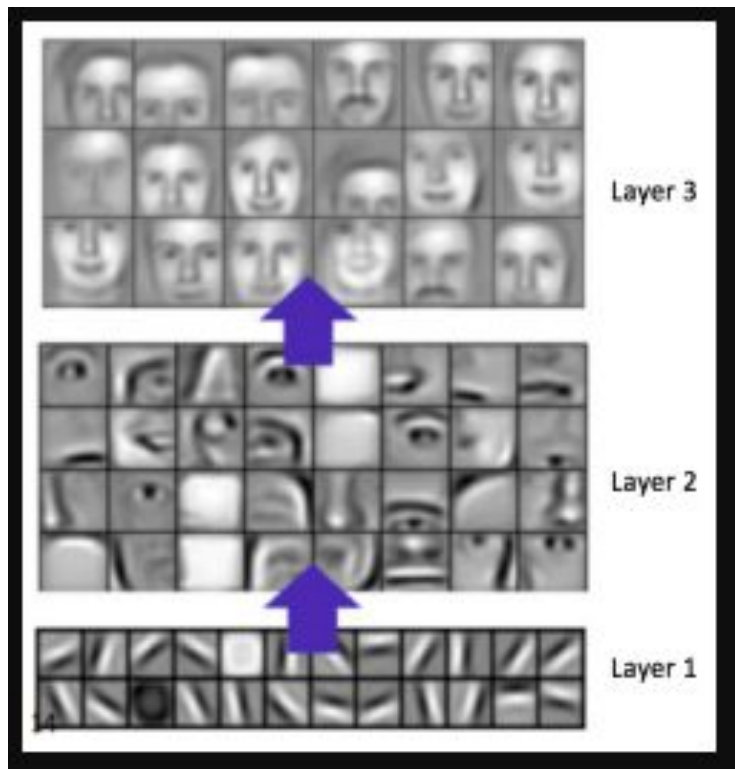
Keras and Tensorflow

Code: MNIST on Dense Network

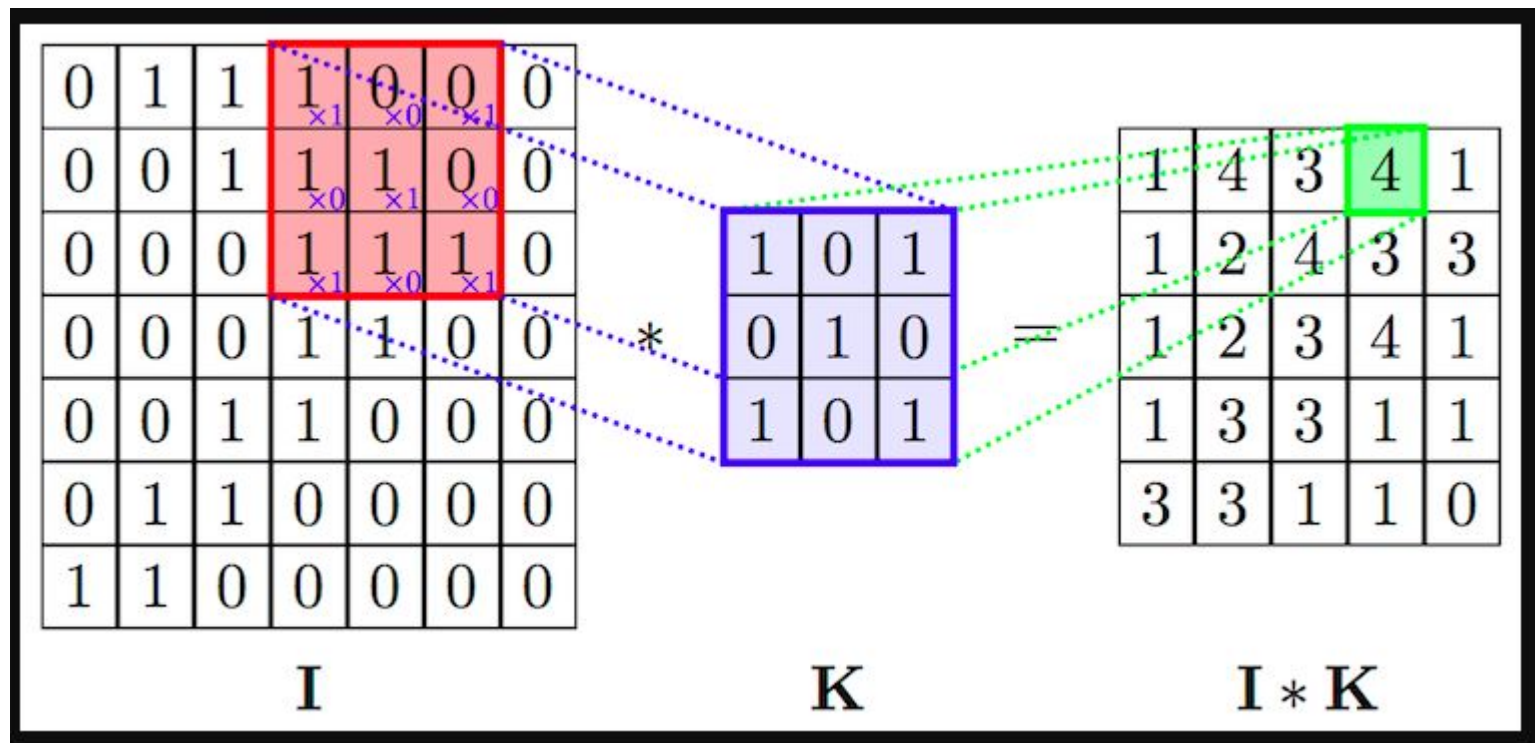
— — —

- <https://keras.io/datasets/>
- <https://github.com/chyld/intro-deep-learning/blob/master/01-fully-connected.ipynb>

Theory of CNNs



Theory of CNNs



Code: Fashion MNIST on CNN

— — —

<https://github.com/chyld/intro-deep-learning/blob/master/02-fashion-cnn-.ipynb>

Deep learning on Tesla GPUs on AWS



Deep Learning AMI (Ubuntu) Version 10.0 - ami-e580c79d

[Select](#)

Free tier eligible

Comes with latest binaries of deep learning frameworks pre-installed in separate virtual environments: MXNet, TensorFlow, Caffe, Caffe2, PyTorch, Keras, Chainer, Theano and CNTK. Fully-configured with NVidia CUDA, cuDNN and NCCL as well as Intel MKL-DNN

64-bit

Root device type: ebs Virtualization type: hvm ENA Enabled: Yes

<u>GPU compute</u>	p2.xlarge	4	61
<u>GPU compute</u>	p2.8xlarge	32	488
<u>GPU compute</u>	p2.16xlarge	64	732
<u>GPU compute</u>	p3.2xlarge	8	61
<u>GPU compute</u>	p3.8xlarge	32	244
<u>GPU compute</u>	p3.16xlarge	64	488

Bonus: Transfer Learning